











`generateStaticParams()` + `dynamicParams` + Full Route Cache

Let's finish this piece properly. This is where **dynamic route $\neq$ dynamic rendering** becomes crystal clear.

We'll use:

```
app/  
└─ products/  
    └─ [id]/  
        └─ page.tsx
```



and:

```
export async function generateStaticParams() {  
  return [  
    { id: "1" },  
    { id: "2" },  
    { id: "3" }  
  ]  
}  
  
export default async function Page({ params }) {  
  const { id } = await params  
  
  return <h1>Product {id}</h1>  
}
```

1 What does `generateStaticParams()` actually do?

Think of `[id]` as a **template**:

```
/products/[id]
```



And:

```
generateStaticParams()
```

returns:

```
id = 1  
id = 2  
id = 3
```



So Next.js can determine:

```
/products/1  
/products/2  
/products/3
```



These are the paths it knows about ahead of time.

Conceptually during build:

generateStaticParams()



```
1  
2  
3
```



render each



```
/products/1  
/products/2  
/products/3
```



Full Route Cache



2 What is actually stored?

Let's pretend:

/products/1



produces:

```
<h1>Product 1</h1>
```

and:

/products/2



produces:

```
<h1>Product 2</h1>
```

and:

/products/3



produces:

```
<h1>Product 3</h1>
```

Then conceptually:

FULL ROUTE CACHE



```
/products/1
  ↓
HTML + RSC Payload
for Product 1

/products/2
  ↓
HTML + RSC Payload
for Product 2

/products/3
  ↓
HTML + RSC Payload
for Product 3
```

Again, the `[id]` component code isn't what gets stored.

The component is executed separately for each parameter value.

3 Now `/products/1`

User requests:

```
/products/1
```



Next.js can find:

Full Route Cache



```
|
└─ /products/1
```

So conceptually:

Browser



```
↓
Next.js
↓
Full Route Cache HIT
```

```
↓  
rendered result  
↓  
HTTP response  
↓  
Browser
```

No need to render the page again.

4 /products/2

Same thing:

```
Browser  
↓  
/products/2  
↓  
Full Route Cache HIT  
↓  
HTML + RSC result  
↓  
Browser
```



5 Now the interesting one: /products/999

We never returned:

```
{ id: "999" }
```

from:

```
generateStaticParams()
```

So:

```
/products/999
```



wasn't pre-rendered during the build.

Now what happens?

It depends on `dynamicParams` .

This is the key.

6 Case A — `dynamicParams = false`

Suppose we add:

```
export const dynamicParams = false
```

So:

```
export async function generateStaticParams() {  
  return [  
    { id: "1" },  
    { id: "2" },  
    { id: "3" }  
  ]  
}  
  
export const dynamicParams = false
```

Now Next.js treats the generated params as the allowed set.

So:

```
/products/1    → ✓  
/products/2    → ✓  
/products/3    → ✓  
/products/999 → ✗
```



The request for `/products/999` results in:

404



Conceptually:

Request `/products/999`

↓

Is 999 in generated params?

↓

NO



```
↓  
dynamicParams = false  
↓  
404
```

7 Case B — `dynamicParams = true`

Now suppose:

```
export const dynamicParams = true
```

(or you're using the default behavior that allows ungenerated params).

We have:

Known at build:

```
1  
2  
3
```



But the user requests:

```
999
```



Next.js can handle that route rather than automatically returning 404.

Conceptually:

```
Request /products/999
```

```
↓  
Not in Full Route Cache  
↓  
Next.js handles route  
↓  
render /products/999  
↓  
response
```



Depending on the route's rendering/caching configuration, that result can subsequently participate in the applicable caching behavior.

The important thing is:

Not being returned by `generateStaticParams()` does not automatically mean "404".

`dynamicParams` controls that behavior.

8 Why is this useful?

Imagine Amazon has:

`/products/[id]`



and:

10 million products



You obviously don't necessarily want:

`next build`

↓

generate 10 million pages



Instead you could pre-render a selected set:

Popular products:



1
2
3
4
5

Then other products can be handled when requested.

Conceptually:

Build

```
|  
├─ /products/1  
├─ /products/2  
├─ /products/3  
└─ ...
```



Later request:

```
/products/999999
```

↓

handled according to `dynamicParams`
and the route's rendering/caching behavior

This is one reason dynamic routes are so powerful.

9 Here's the terminology trap again

We now have:

```
/products/[id]
```



This is a:

Dynamic Route

because `[id]` is a dynamic path segment.

But:

```
/products/1
```



could have been:

Statically generated

because Next.js knew:

```
id = 1
```



from `generateStaticParams()`.

So:

Dynamic Route

≠

Dynamic Rendering



Keep that distinction permanently.

10 Now let's connect this to Full Route Cache

Suppose:


```
generateStaticParams()
```

returns:

```
1  
2  
3
```



Then during build:



That's the big picture.

1 1 Now let's introduce API data

Realistically, your page won't just do:

```
<h1>Product {id}</h1>
```

It'll probably do:

```
export default async function Page({ params }) {
  const { id } = await params

  const response = await fetch(
    `https://api.example.com/products/${id}`
  )

  const product = await response.json()

  return (
    <div>
      <h1>{product.name}</h1>
      <p>{product.price}</p>
    </div>
  )
}
```

Now the pipeline becomes:

```
generateStaticParams()
  ↓
id = 1, 2, 3
  ↓
render /products/1
  ↓
fetch /products/1
  ↓
API data
  ↓
render HTML + RSC
  ↓
Full Route Cache
```



And separately:

```
API response
  ↓
Data Cache
```




if the fetch is configured/cacheable accordingly.

So:

```
API
  |
  ▼
Data Cache
  |
  ▼
```





12 One subtle thing about `generateStaticParams()`

This function is primarily about:

Which dynamic path parameters should be generated ahead of time?

It is **not** a general "make this route static" switch.

For example:

```
generateStaticParams()
```

could give:

```
1  
2  
3
```



but if your page also does:

```
cookies()
```

you've introduced request-specific information.

So don't think:

```
generateStaticParams()  
↓  
"STATIC!!!"
```



Instead:


```
generateStaticParams()  
↓
```



"These parameter values are known ahead of time."

Then Next.js still has to consider the route's other rendering dependencies.

🎯 Our Layer-1 decision tree is getting complete



🔥 Where we are now

We've covered:

- ✓ Static route
- ✓ Dynamic route
- ✓ Cacheable fetch
- ✓ no-store
- ✓ SSG
- ✓ SSR
- ✓ ISR
- ✓ params
- ✓ searchParams
- ✓ generateStaticParams()
- ✓ dynamicParams
- ✓ Full Route Cache



The **last major pieces of Layer 1** are:

```
cookies()  
headers()
```



and then we need to connect **all of them into one actual request/rendering decision tree**.

After that, we'll do a **big end-to-end exercise** where I give you several real Next.js pages and you tell me:

```
Static or Dynamic?  
SSG / SSR / ISR?  
Data Cache?  
Full Route Cache?  
What happens at request time?
```



Then I'll correct your reasoning. That's the point where this stuff should really stick. 🧠



ok



Layer 1 — cookies() and headers() Deep Dive

Now let's finish the two big **request-time APIs**.

We've already seen the rule:

```
cookies()
headers()
  ↓
request-specific information
  ↓
traditionally Dynamic Rendering
```



But I want you to understand **why**, rather than memorize it.

1 Start with a static page

</> TypeScript



```
export default function Page() {
  return <h1>Products</h1>
}
```

Every request gets exactly the same result:

```
Request A → Products
Request B → Products
Request C → Products
```



So Next.js can do:

```
BUILD
  ↓
render once
  ↓
Full Route Cache
  ↓
reuse for requests
```



2 Now introduce cookies()

</> TypeScript



```
import { cookies } from "next/headers"

export default async function Page() {
  const cookieStore = await cookies()

  const theme = cookieStore.get("theme")?.value
}
```

```
return <h1>Theme: {theme}</h1>
}
```

Now imagine:

Request A

Cookie: theme=dark



produces:

HTML



```
<h1>Theme: dark</h1>
```

Request B

Cookie: theme=light



produces:

HTML



```
<h1>Theme: light</h1>
```

So ask the fundamental question:

Can Next.js generate **one HTML result at build time** that is correct for both users?

Can one result work?



NO



Because the result depends on the request.

3 Where does the cookie actually come from?

This is important.

The browser sends a request:

```
Browser
  |
  | GET /products
  | Cookie: theme=dark
  ▼
Next.js
```

The cookie is part of the **incoming HTTP request**.

Then:

```
</> TypeScript
cookies()
```

reads that request information.

So:

```
Browser
  |
  | Cookie header
  ▼
HTTP Request
  |
  ▼
Next.js
  |
  ▼
cookies()
  |
  ▼
theme = "dark"
  |
  ▼
render
```

This is why it is inherently request-dependent.

4 Compare with normal data

Consider:

```
</> TypeScript
const products = await fetch(
  "https://api.example.com/products"
```

```
)
```

If that fetch is cacheable, Next.js can potentially say:

```
"I can reuse this data."
```



But with:

```
</> TypeScript
```

```
cookies()
```



the value comes from:

```
THE CURRENT REQUEST
```



There isn't one universal cookie value.

5 That's why Full Route Cache becomes a problem

Suppose Next.js tried to put:

```
HTML
```

```
</>
```



```
<h1>Theme: dark</h1>
```

into the Full Route Cache.

Then:

```
User B
```

```
Cookie: theme=light
```



comes along.

If Next.js reuses that cached route:

```
HTML
```

```
</>
```



```
<h1>Theme: dark</h1>
```

✗ Wrong.

So the traditional model says:

```
cookies()
  ↓
request-specific
  ↓
dynamic rendering
  ↓
route isn't reusable as one static result
```



6 headers() works exactly the same way

Example:

```
</> TypeScript

import { headers } from "next/headers"

export default async function Page() {
  const headersList = await headers()

  const language =
    headersList.get("accept-language")

  return <h1>Language: {language}</h1>
}
```



The browser sends something like:

```
GET /products
Accept-Language: en-US
```



Then Next.js:

```
Request
  ↓
headers()
  ↓
"en-US"
  ↓
render
```



Another request might send:

```
Accept-Language: fr-FR
```



So the rendered result could differ.

7 Here's the really useful distinction

`cookies()` and `headers()` don't make the **URL** dynamic.

Suppose:

```
app/products/page.tsx
```



URL:

```
/products
```



This is still a static **route structure**.

But:

```
</> TypeScript
```



```
cookies()
```

makes its rendering request-dependent under our Layer-1 model.

So:

```
/products
```



```
├── Route structure: STATIC
└── Rendering: DYNAMIC
```

This is exactly the distinction you asked about earlier.

8 Now let's combine `[id]` with `cookies()`

Suppose:

```
app/products/[id]/page.tsx
```



and:

```
</> TypeScript
```



```
export default async function Page({ params }) {
  const { id } = await params
```

```
const cookieStore = await cookies()
const theme = cookieStore.get("theme")?.value

return (
  <div>
    <h1>Product {id}</h1>
    <p>Theme: {theme}</p>
  </div>
)
```

Now we have:

Dynamic route

```
/products/[id]
```



AND:

Dynamic rendering dependency

```
cookies()
```



So:

```
Dynamic URL
  +
request-specific data
  ↓
traditionally Dynamic Rendering
```



9 What about headers() that doesn't affect UI?

Here's a subtle question.

Suppose:

```
</> TypeScript

const headersList = await headers()

const userAgent = headersList.get("user-agent")

return <h1>Hello</h1>
```



The `userAgent` isn't actually used in the rendered output.

You're probably thinking:

"Then why should the page become dynamic? The HTML is still always `<h1>Hello</h1>`."

Excellent question. 🙌

The **traditional Next.js mental model** treats use of request-time APIs as a signal that the route has a request-time dependency.

But this is exactly the kind of situation where **modern Cache Components changes the architecture**.

In Layer 2, we can isolate request-time work rather than thinking:

```
request API anywhere
      ↓
entire route dynamic
```



So don't over-generalize the Layer-1 rule.

🔥 100 cookies() can also be used for reading AND modifying cookies

Reading:

📄 TypeScript

```
const cookieStore = await cookies()

const token = cookieStore.get("token")
```



is request-time information.

But cookie mutation is a separate concern:

📄 TypeScript

```
cookieStore.set(...)
cookieStore.delete(...)
```



Those operations belong to mutation/request handling contexts such as Server Actions or Route Handlers, rather than ordinary static rendering.


We'll cover that later when we revisit:

Server Actions
Route Handlers
cookies
revalidation



1 1 Now let's connect everything

We've now got:




And:

fetch(no-store)
↓
request-time data
↓
DYNAMIC (traditional)



🧠 1 2 The complete Layer-1 decision tree

Here's the one diagram I want you to remember:



Typical Layer-1 signals:

Usually compatible with static rendering

- ✓ ordinary computation
- ✓ cacheable fetch
- ✓ revalidated fetch
- ✓ `generateStaticParams()`

Traditionally pushes toward dynamic rendering

- ⚠ `fetch({ cache: "no-store" })`
- ⚠ `cookies()`
- ⚠ `headers()`
- ⚠ request-dependent values

🎯 One last thing before we leave Layer 1

We now know the **individual rules**.

But knowing rules isn't enough.

The next lesson should be a **real-world debugging exercise**.

I'll give you pages like:

Example A

```
</> TypeScript

const products = await fetch(url)

const cookieStore = await cookies()

return ...
```



Example B

```
</> TypeScript

const products = await fetch(url, {
  next: { revalidate: 300 }
})

return ...
```



Example C

```
</> TypeScript

const data = await fetch(url, {
  cache: "no-store"
})

const headersList = await headers()

return ...
```



Example D

```
/products/[id]
+
generateStaticParams()
+
revalidate
```



And you'll determine:

1. Dynamic route or static route?
2. Static or dynamic rendering?
3. SSG / SSR / ISR?
4. Data Cache?
5. Full Route Cache?
6. What happens during build?
7. What happens during request?
8. What reaches the browser?
9. What can the CDN cache?



Once you can reason through those without memorizing a table, **Layer 1 is done.** 🔥

Then we'll jump into **Layer 2 — Cache Components** / `use cache` / **Suspense** / **PPR**, where the old mental model gets upgraded.





Ahhh, yes — I see the exact confusion now. 😊 You're asking:

"If `/products` is cached and doesn't depend on the cookie, and I change a cookie in the browser, how does the server ever learn that the cookie changed?"

The answer is: **it learns it on the next request that actually reaches the server.** And if that request is satisfied entirely by a CDN cache, the origin server may **not need to know at all.**

Let's separate these cases.

1. Cookie changes in the browser

Suppose:

```
Browser
_____
theme = light
```



Server Action changes it:

```
cookies().set("theme", "dark")
```



Response:

```
</> http
Set-Cookie: theme=dark
```



Browser now has:

Browser

theme = dark



At this point:

Server: ❌ hasn't received "theme=dark" as a new request

Browser: ✅ knows theme=dark



2. Then user requests `/products`

Browser makes:

```
</> http
```

```
GET /products
```

```
Cookie: theme=dark
```



Now there are two possibilities.

Case A — CDN handles it

Browser

| GET /products

| Cookie: theme=dark



CDN

| Cache HIT



Cached /products response



The request **doesn't reach Next.js origin**.

Therefore:

Next.js server

❌ doesn't learn about this request

❌ doesn't need to



The CDN has already answered the request.

And that's perfectly fine if `/products` doesn't care about `theme`.

3. So when does Next.js actually learn the new cookie?

When there is a request that reaches Next.js.

For example:

```
graph TD; Browser -- "Cookie: theme=dark" --> CDN; CDN -- "Cache MISS" --> Nextjs[Next.js]; Nextjs --> cookies[cookies()];
```

Now Next.js can see:

theme=dark

So:

```
graph TD; A[Cookie changed] --> B[Browser stores new cookie]; B --> C[Future HTTP request]; C --> D[Cookie automatically attached]; D --> E[If request reaches server]; E --> F[Server can read new cookie];
```

4. But here's the important realization 💡

You don't actually need to "tell the server that the cookie changed" as a separate operation.

The cookie itself is the state.

The browser effectively says:

"From now on, whenever I make an applicable request, I'll send `theme=dark`."

So imagine:

10:00

Browser:
theme=light



Server Action:

Set-Cookie: theme=dark



Browser:

theme=dark



Then:

10:05
GET /profile
Cookie: theme=dark



Server sees it.

There was no separate:

UPDATE_COOKIE /api/notify-server



required.

5. What if `/products` is completely cached?

Suppose:

`/products`



doesn't care about cookies.

CDN has:

`/products` → Product HTML



User requests it:

```
Browser
Cookie: theme=dark
  |
  ▼
CDN
  |
  ▼
Cached Product HTML
```



Next.js doesn't learn about the cookie.

And that's okay.

Because the page doesn't use it.

It's exactly like sending someone a letter that says:

```
"Here is my favorite color: dark"
```



but the person only needs to answer:

```
"What products are available?"
```



They don't need to process your favorite color. 😊

6. Now compare with `/profile`

Suppose `/profile` says:

```
</> TypeScript

const cookieStore = await cookies()
const theme = cookieStore.get("theme")?.value
```



Now the cookie matters.

Browser:

```
</> http

GET /profile
Cookie: theme=dark
```



If CDN cannot safely serve a generic cached response:

```
Browser
  ↓
CDN
  ↓
Next.js
  ↓
cookies()
  ↓
theme = dark
  ↓
render
```



Now Next.js **does** learn about it.

7. There's one more scenario: you WANT the server to know immediately

Suppose you have some client-side cookie that you change using JavaScript:

```
</> JavaScript
```

```
document.cookie = "theme=dark"
```



The server isn't immediately notified.

If you specifically need the server to know immediately, you make a request:

```
</> JavaScript
```

```
document.cookie = "theme=dark"
```

```
await fetch("/api/theme", {
  method: "POST"
})
```



Now:

```
Browser
|
| Cookie already changed
| Cookie: theme=dark
▼
POST /api/theme
▼
```



Next.js



Server knows theme=dark

But normally you don't need this extra request unless your application has a reason to synchronize server-side state.



This is the BIG distinction

There are actually **three events**:

Event 1 — Cookie changes

Server → Set-Cookie → Browser



or JavaScript changes it.

Event 2 — Browser makes a request

Browser → HTTP Request + Cookie → network



Event 3 — Somebody processes that request

Could be:

Browser



CDN



response



or

Browser



CDN



Next.js



If CDN answers it:

Next.js doesn't see it.



If CDN forwards it:


Next.js sees the cookie.



🔥 The sentence that should make this click

A cookie is not a live connection between the browser and server. It's client-side state that gets included in applicable HTTP requests.

So:



In **Request A**, Next.js never knows the cookie existed.

In **Request B**, Next.js sees the cookie.

And that's completely normal. 👍

The next thing that will make this whole topic click is understanding **what exactly a CDN cache key is and whether cookies are part of that key**—because that determines whether two users with different cookies can receive the same cached response.



ok

Perfect 😊 Let's lock this piece in.


🧠 Next: CDN Cache Key + Cookies

The question we'll answer is:


If two users request `/products` but have different cookies, how does the CDN know whether it can give them the same cached response?

We'll trace:

User A
Cookie: theme=dark



User B
Cookie: theme=light



Then we'll understand:

1. What a cache key actually is
2. Why `/products` normally maps to one cached response
3. How `Vary` works
4. Whether cookies can be part of cache variation
5. Why personalized pages generally shouldn't be blindly CDN-cached
6. How this differs from **Next.js Data Cache / Full Route Cache / Router Cache**
7. Finally, we'll return to our **Layer 1 static vs dynamic exercise**. 🚀



ok

Alright 😊 let's do it. This is one of those concepts that makes **CDN caching + cookies** suddenly click.

🌐 **CDN Cache Key — What actually identifies a cached response?**

Imagine the CDN receives:

```
</> http
GET /products
```



The CDN needs to answer:

"Do I already have a response for this request?"

It uses something called a **cache key**.

A simplified cache key might be:

```
https://example.com/products
```



So:

```
Request
GET /products
  ↓
Cache Key
example.com/products
  ↓
CDN Cache
  ↓
Cached Response
```



1. Two users request the same URL

Suppose:

User A

```
GET /products
Cookie: theme=dark
```



User B

GET /products
Cookie: theme=light




If the CDN's cache key is simply:

example.com/products



then both requests map to:



So the CDN could potentially return the **same cached response** to both.

That's perfectly fine **if the response doesn't depend on the cookie.**

2. Example where this is perfectly safe

Your `/products` page:

`</>` TypeScript



```
export default async function ProductsPage() {  
  const products = await getProducts()  
  
  return (  
    <ul>  
      {products.map(product => (  
        <li key={product.id}>{product.name}</li>  
      ))}  
    </ul>  
  )  
}
```

The page doesn't care about:

theme
session
language
user ID



So:

```
User A
Cookie: theme=dark
  |
  ▼
/products
  |
  ▼
  CDN
  |
  ▼
Cached products response
```



User B:

```
User B
Cookie: theme=light
  |
  ▼
/products
  |
  ▼
  CDN
  |
  ▼
SAME cached products response
```



That's fine.

Because:

```
Cookie —X—> Products HTML
```



The cookie doesn't influence the response.

3. Now let's make the cookie matter

Suppose:

```
</> TypeScript

export default async function ProductsPage() {
  const cookieStore = await cookies()

  const theme = cookieStore.get("theme")?.value
```



```
return (  
  <div className={theme}>  
    Products  
  </div>  
)  
}
```

Now:

```
Cookie  
  ↓  
theme  
  ↓  
HTML
```



User A:

```
theme=dark  
  ↓  
<div class="dark">
```



User B:

```
theme=light  
  ↓  
<div class="light">
```



So one cached response cannot represent both.

4. Here's the problem with a naive CDN

Suppose CDN already has:

```
Cache:  
  
/products  
  ↓  
<div class="dark">
```



User B arrives:

```
</> http  
  
GET /products  
Cookie: theme=light
```



If CDN says:

"Oh, `/products` is cached! I'll return it."

then:

```
User B
theme=light

↓

CDN

↓

<div class="dark">
```

🔥 Wrong response.

That's a **cache correctness problem**.

5. So how does a CDN handle this?

There are several mechanisms in real systems.

One important HTTP mechanism is:

```
</> http
Vary
```

For example:

```
</> http
Vary: Accept-Language
```

This tells caches:

"The response varies based on this request header."

So conceptually:

```
/products + language=en
/products + language=fr
```

can be treated as different variants.

6. What about cookies?

Here's where things get more nuanced.

You *can* design caching behavior around cookies, but you generally don't want to casually make:

```
Cache Key =  
URL + entire Cookie header
```



because imagine:

```
Cookie:  
session=abc123;  
theme=dark;  
cart=92837;  
...
```



Now every user's session could produce a completely different cache entry.

You'd destroy much of the benefit of shared caching.

You could theoretically get:

```
/products + userA cookie  
/products + userB cookie  
/products + userC cookie  
/products + userD cookie  
...
```



That's basically **personalized caching**.

7. This is why personalized pages are tricky

Imagine:

```
<> TypeScript  
  
const session = cookies().get("session")  
  
return <h1>Welcome, {sessionUser.name}</h1>
```



User A:

Welcome Gopi



User B:

Welcome Rahul



You absolutely cannot have:

/products



Welcome Gopi



as one generic shared CDN response.

Otherwise:

User B



CDN



"Welcome Gopi"



😞 That's obviously terrible.

So systems generally avoid treating such personalized responses as ordinary shared public cache entries, or explicitly vary/cache them according to an appropriate strategy.

8. Now connect this to Next.js

This is why our earlier statement was:

cookies()



request-dependent



dynamic rendering



Under the **traditional Layer 1 model**, Next.js doesn't say:

"I'll just put this personalized page in the generic Full Route Cache."

Instead, the route is treated as request-dependent.

Conceptually:

```
Request
  ↓
Next.js
  ↓
cookies()
  ↓
render THIS request
  ↓
response
```



9. But here's a very important distinction 🚨

Don't think:

"If a request contains a cookie, CDN can never cache it."

That's too simplistic.

The correct idea is:

The presence of a cookie doesn't automatically make every response uncacheable. What matters is whether the response is safe to reuse for another request.

For example:

Cookie: theme=dark



but `/products` doesn't use theme.

The response could still potentially be shared.

Whereas:

Cookie: userId=123



and the response says:

Welcome Gopi



That's personalized.


Different situation.

10. Now our three cache layers become clearer

Let's put them side-by-side.

Layer	What it caches	Where
Data Cache	Fetches data	Next.js/server infrastructure
Full Route Cache	Rendered route result	Next.js/server infrastructure
CDN/HTTP cache	HTTP responses	CDN/edge/intermediary
Router Cache	RSC payload	Browser/client

So:



And importantly, **these aren't duplicate copies of exactly the same thing for the same purpose.**

11. One killer example 🔥

Imagine your page:

</> TypeScript



```
export default async function Page() {  
  const cookieStore = await cookies()  
  const theme = cookieStore.get("theme")?.value  
  
  const response = await fetch(  
    "https://api.example.com/products",  
    {  
      next: { revalidate: 60 }  
    }  
  )  
  
  const products = await response.json()  
  
  return (  
    <div className={theme}>  
      ...  
    </div>  
  )  
}
```

Now there are **two different kinds of information**:

Cookie

theme=dark



Request-specific.

Products API data


iPhone
MacBook
AirPods




Potentially cacheable/reusable.

So conceptually:

Cookie






This is a really important insight:

One route can have request-specific information AND separately cacheable data dependencies.


And this becomes **hugely important when we get to Layer 2 / Cache Components**, because modern Next.js gives you more precise control over which parts are cached and which parts remain request-dependent.

Your mental model now

Don't think:



Instead:



That's the whole game. 🎯

And now we're ready for the **Layer 1 debugging exercise**. I'll give you a route and code, and you tell me what happens. I'll deliberately mix **dynamic URL segments**, **cookies**, `searchParams`, `fetch`, `no-store`, **and** `revalidate` so you have to distinguish all of them.



ChatGPT can make mistakes. Check important info.