



Almost! 🤔 But based on **what we've actually completed**, here's the clean status:

🏁 Current status

- ✅ 6.11 — RSC / Flight / RSC Payload
- ✅ 6.12 — Serialization & References
- ✅ 6.13 — Server ↔ Client boundaries
- ✅ 6.14 — `router.refresh()` vs navigation vs reload
- ✅ 6.15 — Router Cache + RSC navigation
- ⌚ 6.16 — DevTools hands-on experiment

Show more ▾



Yep 🤔 — we're starting 6.17 now.


🚀 Next.js 6.17 — Final Rendering + Caching Mental Model


This is the “**connect every dot**” lesson. Instead of learning another isolated Next.js feature, we're going to reason through what actually happens from:

user request → server → rendering → caches → RSC payload → browser → navigation

6.17.1 — The complete picture

Think of a modern Next.js App Router application as this pipeline:





But there is an important distinction:

HTML and RSC Payload have different jobs.

6.17.2 — What does the server actually send?

Suppose we have:

</> TypeScript

`app/products/page.tsx`



</> TypeScript

```

export default async function ProductsPage() {
  const products = await getProducts();

  return (

```



```

    <div>
      <h1>Products</h1>

      <ProductList products={products} />

      <AddToCartButton />
    </div>
  );
}

```

Imagine:

</> TypeScript



ProductList = Server Component
AddToCartButton = Client Component

The server can execute:

```

ProductsPage
  ↓
getProducts()
  ↓
ProductList

```



and produce the information necessary to render the UI.

The initial request can result in:

1. HTML

Something like:

HTML



```

<div>
  <h1>Products</h1>
  <ul>
    <li>iPhone</li>
    <li>MacBook</li>
  </ul>
  ...
</div>

```

This allows the browser to **show something quickly**.

2. RSC Payload

React also receives the **React Server Component representation** needed to understand the Server Component tree and Client Component references.

Conceptually:

```
Server Component tree
  ↓
React Flight / RSC Payload
  ↓
Browser React runtime
  ↓
Reconstruct/merge component tree
```



The actual Flight format is an implementation detail and isn't simply "HTML in another format."

6.17.3 — Why do we need RSC Payload if HTML already exists?

This is one of the **most important questions** you've been asking throughout this block.



Because:

HTML tells the browser what to display.

RSC Payload tells React about the Server Component result/tree and how it relates to the client-side React application.

Think:

```
HTML
= initial visual document
```

```
RSC Payload
= React's server-generated component information
```



So they aren't competing representations.

They're complementary.

6.17.4 — Why doesn't the browser receive Server Component JavaScript?

Suppose:

</> TypeScript



```
export default function ProductList() {  
  return <div>Products...</div>;  
}
```

is a Server Component.

Its source code isn't shipped to the browser simply so the browser can execute it.

Instead:

Server



ProductList()

↓

executed on server

↓

result represented in RSC Payload

↓

HTML / RSC information sent

↓

Browser

React receives result

↓

updates UI

That's the core benefit of React Server Components.

The browser doesn't need to download and execute every Server Component.

6.17.5 — Then what JavaScript DOES reach the browser?

Client Components.

For example:

</> TypeScript



```
"use client";  
  
export function AddToCartButton() {  
  const [count, setCount] = useState(0);  
  
  return (  
    <button onClick={() => setCount(count + 1)}>
```

```
Add {count}
</button>
);
}
```

This component requires browser-side JavaScript because:

```
useState
onClick
event handling
browser interaction
```



So:

```
Server Component
  |
  | server-only
  ▼
Server
Client Component
  |
  | JS bundle
  ▼
Browser
```



And this is where **serialization** and **Client References** become important.

6.17.6 — Serialization: the boundary

Remember our earlier rule:

```
Server Component
  |
  | props
  ▼
Client Component
```



The data crossing this boundary needs to be representable in the RSC/React protocol.


For example:

```
</> TypeScript
<ClientComponent
  name="Gopi"
```



```
count={10}  
/>
```

Conceptually:



But you can't casually pass arbitrary server-only things such as:

```
</> TypeScript  
  
<ClientComponent databaseConnection={db} />
```

because that object isn't something you can simply serialize and transfer to the browser.

This is why **Server/Client boundaries + serialization + references** are one connected concept.

6.17.7 — Now let's add caching

Here's where Next.js can feel like a bowl of spaghetti until you separate the layers. 😂

There isn't just **"the Next.js cache."**

We learned several different mechanisms.

Data Cache

Caches data resulting from supported server-side data fetching.

Conceptually:

```
fetch(...)  
↓  
Data Cache
```

Full Route Cache

Historically, this represented cached rendered output for routes.

With the newer Cache Components architecture, the mental model becomes more nuanced: cached portions/data can be persisted and reused rather than thinking of the entire route as one giant immutable blob.

Router Cache

This is on the **client**.

Browser



Router Cache



It helps Next.js reuse previously received RSC data during client-side navigation.

Request Memoization

This isn't simply another persistent cache.

It deduplicates certain identical requests during a render/request lifecycle.

6.17.8 — The BIG distinction: server caches vs browser cache

This is worth tattooing on your brain. 😊

SERVER



Data Cache
Render/cache layers

RSC / HTML

BROWSER

Router Cache
Browser state

So when you say:

"The page is cached"

you should immediately ask:

Which cache?

Because:

Data Cache

≠

Router Cache

≠

HTTP/CDN cache

≠

browser HTTP cache



6.17.9 — Initial page load vs navigation

This is another huge distinction.

Initial request

User enters:

/products



The browser needs the initial page.

Conceptually:

Browser

GET /products



Next.js server

— routing

— Server Components

— data fetching

— caching

— rendering


→ HTML

→ RSC information

+

JS bundles





6.17.10 — Now click `<Link href="/about">`

This is where things get interesting.

You're already inside a Next.js application.

So instead of doing a traditional full document navigation:



```
GET /about
  ↓
new HTML document
  ↓
replace current document
```

A diagram showing a vertical flow: 'GET /about' followed by a downward arrow to 'new HTML document', followed by another downward arrow to 'replace current document'.

Next.js can perform a **client-side navigation**.

Conceptually:



```
Browser
  |
  | navigation to /about
  ↓
Next.js server
  |
  | RSC / Flight response
  ↓
Browser
  |
  ↓
React Router / App Router
  |
  ↓
merge new RSC result
  |
  ↓
update affected UI
```

A diagram showing a vertical flow: 'Browser' followed by a line with 'navigation to /about' and a downward arrow to 'Next.js server'. From 'Next.js server', a line with 'RSC / Flight response' and a downward arrow leads to another 'Browser'. From this 'Browser', a downward arrow leads to 'React Router / App Router'. From there, a downward arrow leads to 'merge new RSC result', and finally a downward arrow leads to 'update affected UI'.

The existing document isn't necessarily thrown away.

That's why client-side navigation feels much more like:

"change this part of my application"



rather than:

"load a completely new website"



6.17.11 — Does navigation fetch HTML again?

This was one of your earlier questions.

For a normal client-side App Router navigation:

the important response is the RSC/Flight data used to update the route.

It isn't the same as doing a traditional browser document request where the browser receives a complete new HTML document and replaces the page.

So:

Initial load



HTML + RSC/Flight + JS

versus:

Client navigation



RSC/Flight data

↓

React updates existing application

That distinction is **extremely important for interviews**.

6.17.12 — What about a completely different page structure?

Suppose:

/products



looks like:

Header
ProductList
Footer



and you navigate to:

/dashboard



which looks like:

DashboardSidebar
DashboardCharts
DashboardTable



You might wonder:

"But the HTML structure changes. Don't we need new HTML?"

The answer is:

The browser's DOM can absolutely change.

But that doesn't mean Next.js needs to replace the entire HTML document.

Instead:

Current React tree



Header
ProductList
Footer

gets updated/reconciled with the new route's React tree:

DashboardSidebar
DashboardCharts
DashboardTable



React manipulates the DOM to make the UI match the new tree.

So:

HTML document replacement ❌



DOM updates/reconciliation 

That's the subtle distinction.

6.17.13 — Where does Router Cache fit?

Imagine you visited:

/products



Then:

/about




Then:

/products



The client-side Router Cache may allow Next.js to reuse previously received route data rather than fetching everything from scratch again, depending on the current caching/revalidation behavior.

Conceptually:



So navigation can become:

```
<Link>
  ↓
check/reuse router state/cache where applicable
  ↓
fetch missing/stale RSC data if needed
  ↓
React updates tree
```



6.17.14 — `router.refresh()`

Now compare that with:

`</> TypeScript`

```
router.refresh();
```



This does **not** mean:

```
delete browser page  
reload entire website
```



Instead, conceptually:

```
Current route  
  |  
  ▼  
request fresh Server Component result  
  |  
  ▼  
receive RSC/Flight response  
  |  
  ▼  
React merges it into current UI
```



So:

Browser reload

```
Reload document  
↓  
new page request  
↓  
HTML again  
↓  
new document
```



Client navigation

```
Navigate  
↓  
RSC/Flight  
↓  
new route tree  
↓  
React updates UI
```



router.refresh()

Same route

↓

request fresh Server Component result

↓

RSC/Flight

↓

React updates current UI

🔥 That's the three-way distinction you want to be able to explain without hesitation.


6.17.15 — Now put rendering into the picture

Rendering isn't simply:

"SSR or CSR?"

Modern Next.js is more nuanced.

Think:



And the actual behavior depends on things such as:

- dynamic request APIs

- uncached data
- cached data
- explicit caching
- Cache Components
- route configuration
- Suspense boundaries
- streaming

So don't reduce modern Next.js to:

```
SSR = server
SSG = static
CSR = browser
```



That's too simplistic.

6.17.16 — Streaming

Suppose your page has:

```
Header          → fast
ProductList     → fast
Recommendations → slow
Reviews         → VERY slow
```



Without streaming, the user may have to wait for everything before seeing the final result.

With Suspense/streaming:



So the browser can progressively receive/render pieces.

This is one reason:

Rendering
+
RSC
+
Suspense
+
Streaming



need to be understood together.

6.17.17 — The ultimate mental model 🧠

If an interviewer gives you a random Next.js scenario, walk through these questions in **this order**:

① What route?

Which segment?
Dynamic?
Nested?
Layout?



② Server or Client?

Is this component Server?
Does it need "use client"?



③ What data?

fetch?
database?
params?
cookies?
headers?



④ Is the data cached?

Cached?
Uncached?



Revalidated?



⑤ Does the request require dynamic behavior?

cookies()
headers()
uncached request?
other request-dependent data?



⑥ How does it render?

Static/cacheable?
Dynamic/request-time?
Streaming?



⑦ What gets sent?

Initial request:

HTML
+
RSC/Flight
+
required JS



Navigation:

RSC/Flight
+
required client-side code



⑧ What happens in the browser?

React receives data
↓
reconciles tree
↓
updates DOM



↓
Client Components hydrate/work

↓

⑨ What happens on the next navigation?


Router Cache?
Fresh RSC request?
Existing layouts reused?




That is your **Next.js** reasoning framework.

6.17.18 — The one diagram I want you to remember

Save this one mentally:





If you can **walk through that diagram verbally**, you understand the architecture rather than just memorizing terminology. 🔥

🎯 6.17.19 — Interview test

Now let's test you.

Don't look anything up. Answer these in your own words:

Q1

A user directly enters:

/products/123



in the browser address bar.

Walk me through what happens from request → server → browser.

Q2

The user is already on:

/products



and clicks:

```
</> TypeScript
```



```
<Link href="/products/123">
```

How is this different from Q1?

Q3

A Server Component fetches:

```
</> TypeScript
```



```
const product = await fetch("/api/product/123");
```

and a Client Component receives:

```
</> TypeScript
```



```
<ProductButton product={product} />
```

Explain:

What crosses the Server → Client boundary?

Q4

What is the difference between:

```
Data Cache  
Router Cache
```



and where does each one live?

Q5

Explain the difference between:

```
browser reload  
navigation  
router.refresh()
```



Q6 — the tricky one 🐱

A page contains:

```
<Header />

<Suspense fallback={<Loading />}>
  <SlowProductList />
</Suspense>

<AddToCartButton />
```



Explain what can happen during rendering and why **streaming + RSC + Client Components** can all participate in the same page.

Q7 — final boss 🐉

Someone says:

"Server Components are rendered into HTML, so RSC Payload is unnecessary."

Do you agree? Explain **exactly why or why not**.

Answer these **7 without worrying about perfect wording**. I'll grade your reasoning, correct the gaps, and then we'll officially mark **6.17 🏁 COMPLETE**.



Absolutely. Let's pause the interview test for a second because **Request Memoization** is one of those Next.js terms that sounds like caching but is actually quite different. 🤔

And your **Q1 is mostly right** — I'll correct one important detail after we nail memoization.

1. Request Memoization — what exactly is it?

The easiest way to understand it:

Request Memoization prevents the same `fetch()` request from being performed multiple times during the same server render.

It is basically **deduplication**, not persistent caching.

Imagine this Server Component tree:

```
ProductPage
├── ProductDetails
│   └── fetch("/api/products/123")
├── ProductReviews
│   └── fetch("/api/products/123")
└── RelatedProducts
    └── fetch("/api/products/123")
```

Three components ask for exactly the same data.

Without deduplication, conceptually:

```
ProductDetails
  ↓
fetch("/api/products/123") → Server/API


ProductReviews
  ↓
fetch("/api/products/123") → Server/API

RelatedProducts
  ↓
fetch("/api/products/123") → Server/API
```

That's unnecessary.

With request memoization:

```
fetch("/api/products/123")
```



The first matching request establishes the result, and subsequent matching requests can reuse that result **within the relevant render/request lifecycle**.

2. Memoization ≠ Data Cache

This distinction is 🔥 important.

Request Memoization

Think:

"Have I already requested this exact thing during this render?"

Data Cache

Think:

"Do I have a stored result that I can reuse across requests/renders?"

For example:

REQUEST 1

`fetch("/products/123")`

↓

request memoization

↓

one request

REQUEST 2

new request/render

↓

request memoization starts fresh

but...

Data Cache may already contain
`/products/123`



So:

Request Memo



Request Memoization

This isn't simply another persistent cache.

It deduplicates certain identical requests during a render/request lifecycle.-explain in detail???

1) So, for the first question, if a user directly enters slash product slash 123 in the browser address bar, so first what happens? The request goes to the server, and there, according to whatever the code we have written, the HTML plus server-side components, RSI payload, plus client JS components come to the browser. And

Show more ▾

Yes! 😊 Your overall mental model is **quite solid now**. The biggest things to fix are:

1. **Request Memoization** — you haven't answered this yet.
2. **Q2** — Router Cache doesn't necessarily mean the requested route is already cached.
3. **Q3** — you're mixing up *data* crossing the boundary with *references*.
4. **Q6** — Suspense/streaming and hydration need a little tightening.
5. **Q7** — your core intuition is right, but the exact role of RSC Payload is slightly different from "where to place the Client Component."

Let's go one by one.

🔥 First: Request Memoization in DETAIL

This is the part you specifically asked about.

Imagine this component tree

</> TypeScript



```
export default async function ProductPage() {  
  return (  
    <>  
      <ProductDetails />  
      <ProductReviews />  
    </>  
  );  
}
```

And:

```
</> TypeScript   
  
async function ProductDetails() {  
  const product = await fetch("/api/product/123");  
  ...  
}
```

and:

```
</> TypeScript   
  
async function ProductReviews() {  
  const product = await fetch("/api/product/123");  
  ...  
}
```

Both components make:

```
</> TypeScript   
  
fetch("/api/product/123")
```

during the same render.


You might initially think:

```
ProductDetails   
  ↓  
fetch()  
  ↓  
API  
  
ProductReviews  
  ↓  
fetch()  
  ↓  
API
```

So two requests.

But Next.js/React can **deduplicate identical** `fetch` requests during the render.

Conceptually:



Both components can receive the same result.

Why is this called "memoization"?

Memoization generally means:

Remember the result of something so that if the same computation is requested again, you don't have to perform it again.

Simple JavaScript example:

`</>` JavaScript

```
function expensiveCalculation(x) {
  // expensive work
  return x * 10;
}
```

If you repeatedly ask:

`</>` JavaScript

```
expensiveCalculation(5)
expensiveCalculation(5)
```

```
expensiveCalculation(5)
```

memoization can conceptually turn that into:

```
First 5 → calculate → remember result  
Second 5 → reuse result  
Third 5 → reuse result
```



Request memoization applies the same basic idea to certain requests.

⚠ But here's the REALLY important part

Request Memoization is **not** saying:

"Store this API response forever."

That's where it differs from a persistent cache.

Think:

Request Memoization

ONE RENDER / REQUEST



```
fetch(A)  
  ↓  
perform request  
  ↓  
remember result for this render  
  ↓  
fetch(A) again  
  ↓  
reuse result
```

When that request/render lifecycle is over, that memoization isn't the thing responsible for serving the next user's request.

Compare this with Data Cache

Suppose:


```
</> TypeScript
```



```
fetch("/api/product/123")
```

is also eligible for persistent Data Cache behavior.

Then you can conceptually have:



So they operate at **different scopes**.



Concrete example

Imagine:

```
</> TypeScript
```

```
export default async function Page() {
  const a = fetch("/api/products");
  const b = fetch("/api/products");

  const [products1, products2] = await Promise.all([a, b]);

  ...
}
```

You're asking for the same resource twice.

Conceptually:

```
fetch("/api/products") ┌  
                        │  
fetch("/api/products") └─> same underlying request
```



The two callers can get the same resulting response.

But now:

</> TypeScript

```
fetch("/api/products?category=shoes")  
fetch("/api/products?category=phones")
```



Those are **different requests**.

Request A
/products?category=shoes

Request B
/products?category=phones



So memoization doesn't magically say:

"They're both products, I'll reuse them."

The request identity matters.



Important modern Next.js nuance

One thing I want you to update in your notes:

Older Next.js explanations often describe **"Request Memoization"** as a Next.js caching layer for GET fetches.

Modern React/Next.js documentation makes the distinction more nuanced: **React extends `fetch` behavior for server rendering and deduplication, while Next.js caching mechanisms such as the Data Cache have their own lifecycle and semantics.**

So for interviews, don't say:

✗ "Request Memoization is a persistent Next.js cache."

Say:

✓ "Request memoization deduplicates identical data requests during a render, so multiple components requesting the same resource don't necessarily cause multiple underlying requests."

That answer is much safer and more accurate.

Now let's grade your answers 🔥

Q1 — Initial request

You said:

request goes to server → HTML + RSC payload + client JS → browser → hydration → UI

✓ **Good.**

But there's one correction:

Don't think:

Server Components → browser



The **Server Component code itself isn't shipped to the browser** as client JavaScript.

Instead:

Server

- executes Server Components
- produces RSC/Flight result
- produces HTML for the initial response
- sends required Client Component JS




Then:

Browser

- displays HTML





So your Q1 is basically:

✅ 8.5/10

Q2 — Navigation

You said:

since RSC payload for `/product/123` exists in Router Cache, we don't need to request server.

⚠️ **This is the important correction.**

You cannot assume:


`/products/123`
↓
already in Router Cache



Just because you're navigating there.

If it isn't cached, Next.js can request the server.

The better model is:



↓
React updates UI

So the key statement is:

Client navigation uses the Router Cache when applicable; otherwise it can request the server for the required RSC payload.

Also, layouts/segments that are already available can be reused while only the necessary new segments are obtained.

Q2: 7.5/10

Your overall navigation model is right; you just assumed the cache hit.

Q3 — Server → Client boundary

This is the one you're struggling with, so let's make it crystal clear.

Suppose:

```
</> TypeScript
// Server Component

const product = await getProduct();

return <BuyButton product={product} />;
```



where:

```
</> TypeScript

"use client";

function BuyButton({ product }) {
  ...
}
```



What crosses?

The data crosses.

For example:

```
Server

product = {
```



```
id: 123,  
name: "iPhone",  
price: 80000  
}
```

↓
Server → Client boundary
↓

Client

BuyButton receives:

```
{  
  id: 123,  
  name: "iPhone",  
  price: 80000  
}
```

That data has to be representable through the React Server Component protocol.

What does NOT cross?

The server's internal implementation doesn't get sent:

```
database connection  
server environment  
API secret  
server-only function implementation
```



Those remain server-side.

And where do References come in?

Suppose:

```
</> TypeScript  
  
<BuyButton />
```



is a Client Component.

The server needs to communicate something like:

"This part of the tree corresponds to this Client Component/module."

That's where **Client References** are important.

Conceptually:

RSC Payload



"Here is a Client Component reference
to module X, with these props."

The browser already has the corresponding Client Component JavaScript
bundle/module.

So:

Server



├ Client Component reference

└ serializable props



RSC Payload



Browser



find/load client module



give it props

That's the missing piece in your explanation.

Q3: 6.5/10

Your intuition was there; you just needed to separate:

DATA

+

CLIENT REFERENCE



from:

SERVER COMPONENT CODE



Q4 — Data Cache vs Router Cache

You said:

Data Cache caches fetched result, Router Cache caches RSC payload, Data Cache server, Router Cache browser.

🎯 **Very good.**

Conceptually:

	Data Cache	Router Cache
Where?	Server-side	Browser/client
What?	Cached data/results	RSC/route segment data
Purpose	Reuse server-fetched data	Make client navigation efficient
Scope	Server	Individual browser session/app

One nuance:

Don't describe Router Cache as simply:

"cache of the page."

It's more useful to think:

client-side cache of previously received RSC/route segment information used by the App Router.

Q4: 9/10

Q5 — Reload vs navigation vs refresh

You said:

Reload

fetch new HTML document

Navigation

request new RSC payload

```
router.refresh()
```

ask server for current page instead of existing cache

That's a **very good answer**.

I'd make one small correction to the last part.

`router.refresh()` doesn't mean:

"ignore every cache."

It requests a fresh RSC response for the current route, but server-side caching can still influence what the server returns.

So:

```
router.refresh()  
  ↓  
new RSC request  
  ↓  
server determines result  
  ↓  
React merges result into current UI
```



not:

```
router.refresh()  
  ↓  
delete every cache
```



Q5: 9/10

Q6 — Streaming

You said:

Header renders, Suspense boundary shows loading, slow product list arrives later, RSC payload reaches React, then UI updates, AddToCart also reaches browser.

Mostly correct. 👍

The important distinction is:

```
<Suspense fallback={<Loading />}>  
  <SlowProductList />  
</Suspense>
```



doesn't mean:

"SlowProductList is a Client Component."

It can absolutely be a **Server Component**.

For example:

```
</> TypeScript

<Suspense fallback={<Loading />}>
  <SlowProductList />
</Suspense>
```

could involve:

```
Server
|
|— render Header
|
|— render AddToCartButton
|
|— wait for SlowProductList
    |
    |— eventually produce RSC/HTML stream
```

The fallback lets the UI make progress while the slow part is pending.

And then the server can progressively stream the result.

Q6: 8/10

Q7 — Your final answer

You said:

Server Components are rendered on server, but RSC payload shows where exactly to place this client component and what's the interconnection between server component and client module.

🟢 **Your core idea is correct.**

But let's make the wording more precise.

RSC Payload isn't primarily:


"Here's the exact DOM position where the Client Component goes."

That's a little too DOM-oriented.

It's better to say:

RSC Payload represents the server-rendered React component tree/result, including information about Client Components and the references needed for React to reconstruct/merge that tree on the client.

So:



Therefore:

✗ "HTML already contains everything, so RSC Payload isn't needed."

is wrong for an App Router RSC application.

HTML is primarily useful for **the initial visual document**.

RSC Payload provides React with the **component-tree information needed for RSC-aware rendering and client-side navigation/updates**.

Q7: 8.5/10



Your overall score

I'd give you roughly:

8.2/10


And honestly, the important thing is *where* your mistakes are.


You're no longer confused about the architecture itself. Your remaining mistakes are mostly **precision issues**:

- ✗ Router Cache always has destination
↓
- ✓ Router Cache may have destination
- ✗ Server Component goes to browser
↓
- ✓ Server executes it; result goes to browser
- ✗ RSC tells browser DOM placement
↓
- ✓ RSC represents React tree/result + Client References + props
- ✗ router.refresh = bypass all cache
↓
- ✓ router.refresh = request fresh RSC result,
while server caching semantics can still apply

And now connect Request Memoization to everything

Here's the final connection:





That's the complete chain. 🔥

And one terminology correction for your notes: you've repeatedly typed "**RSI/RSS payload**" — the term is **RSC Payload**, often described as the **React Flight** payload. 😊



define react server component and what it is exactly to a interviewer

Absolutely. For an interview, I'd keep it **simple first**, then add the technical detail if they ask.

🎯 Interview answer — "What is a React Server Component?"

A **React Server Component** is a React component that is rendered and executed on the server instead of being sent to the browser as JavaScript.

It can directly access server-side resources like databases, files, or server-only APIs, and the result is sent to the client through the **React Server Component**

(RSC/Flight) protocol.

Because its code doesn't need to be shipped to the browser, it can reduce the amount of client-side JavaScript. It also cannot directly handle browser interactions like `onClick` or `useState`; for those, we use a **Client Component** with `"use client"`.

If they ask: "What exactly happens?"

Say:

```
Server Component
  ↓
executes on server
  ↓
produces React tree/result
  ↓
RSC / Flight Payload
  ↓
Browser
  ↓
React uses that information to build/update the UI
```

For example:

```
</> TypeScript

// Server Component
async function ProductPage() {
  const product = await db.products.findById(123);

  return <ProductDetails product={product} />;
}
```

The important point is:

```
db access
  ↓
happens on SERVER
```

The browser doesn't receive the `ProductPage` implementation as client-side JavaScript.

Server Component vs Client Component

Server Component	Client Component
Executes on server	Executes in browser
Can access server-side resources	Runs with browser capabilities
Doesn't ship its component code as client JS	Its JS is shipped to browser
Can be async	Can use browser interaction APIs
Cannot use <code>useState</code> , <code>useEffect</code> , <code>onClick</code> directly	Can use <code>useState</code> , <code>useEffect</code> , <code>onClick</code>
Default in Next.js App Router	Opt in with <code>"use client"</code>

Example:

`</>` TypeScript



```
// Server Component
async function ProductPage() {
  const product = await getProduct();

  return <BuyButton product={product} />;
}
```

`</>` TypeScript



```
// Client Component
"use client";

function BuyButton({ product }) {
  const [loading, setLoading] = useState(false);

  return (
    <button onClick={() => setLoading(true)}>
      Buy {product.name}
    </button>
  );
}
```

So the simplest mental model is:


SERVER

BROWSER



|
|
| Server Component
| executes here

|
|
|



★ One sentence to memorize

"A React Server Component is a component whose rendering logic runs on the server, allowing it to use server-side resources without sending that component's implementation to the browser; the resulting component information is delivered through the RSC/Flight protocol."

That is a **strong interview definition** because it explains not just *where* it runs, but **what** makes it different and **why** it exists.



ChatGPT can make mistakes. Check important info.